

FUNCTION CALL TRACING

XRAY IN LLVM



WHAT TO EXPECT

WHAT, HOW, AND WHEN.

TOPICS TO DISCUSS

- ▶ What is XRay?
- ▶ How do you use XRay?
- ▶ How is XRay implemented?
- ▶ What to expect next?

BEFORE ANYTHING ELSE...

WHY XRAY?

LATENCY & PERFORMANCE DEBUGGING IS HARD.

- ▶ Sampling is an OK approximation for *mean/average* latency, but not for tail latency debugging.
- ▶ Existing solutions require some OS-level support, special privileges, and non-trivial overheads.
- ▶ Manual tracing/logging is not cheap and usually error-prone.

XRAY: HIGH LEVEL FEATURES AND CONSTRAINTS

- ▶ XRay leverages compiler knowledge of code structure to provide instrumentation in interesting places.
- ▶ XRay instrumentation points are cheap, can be deployed in production with little interference/effect.
- ▶ Dynamic control (being able to turn things on/off at runtime) is crucial for production deployments.
- ▶ Gathering data that can be analysed offline.



A WALKTHROUGH

WHAT IS XRAY?

Image from https://upload.wikimedia.org/wikipedia/commons/c/cd/Historical_X-ray_nci-vol-1893-300.jpg


WHAT IS XRAY?

- ▶ Compiler-inserted instrumentation in functions
- ▶ Runtime library for dynamic instrumentation
- ▶ Tools for analysing traces

**COMPILER INSERTED
INSTRUMENTATION.**

NO-OPS IN THE RIGHT PLACES

- ▶ We add "sleds" of instructions that we can overwrite at runtime.
- ▶ We put them in strategic places, to allow us to catch entry and exits at runtime.
- ▶ We put them on "interesting" functions, according to number of instructions.
- ▶ We also put them on explicitly marked functions.



AN EXAMPLE IN C++

CODE.

```
[[clang::xray_always_instrument]] void foo() { printf("Hello, XRay!\n"); }
```

Explicitly annotated functions for "always" or "never" instrumented functions.

clang++ -fxray-instrument ...

```
.globl _Z3foov
.p2align 4, 0x90
.type _Z3foov,@function
_Z3foov:                                     # @_Z3foov
.Lfunc_begin0:
.file 22 "test.cc"
.loc 22 6 0                                # test.cc:6:0
.cfi_startproc
# BB#0:                                     # %entry
.p2align 1, 0x90
.Lxray_sled_0:                               Entry sled.
.ascii "\353\t"
nopw 512(%rax,%rax)
.Ltmp0:
pushq %rbp
.Ltmp1:
.cfi_def_cfa_offset 16
.Ltmp2:
.cfi_offset %rbp, -16
movq %rsp, %rbp
.Ltmp3:
.cfi_def_cfa_register %rbp
subq $16, %rsp
movabsq $.L.str, %rdi
.Ltmp4:
.loc 22 6 48 prologue_end                  # test.cc:6:48
movb $0, %al
callq printf
.loc 22 6 74 is_stmt 0                     # test.cc:6:74
movl %eax, -4(%rbp)                        # 4-byte Spill
addq $16, %rsp
popq %rbp
.p2align 1, 0x90
.Lxray_sled_1:                               Exit sled.
retq
nopw %cs:512(%rax,%rax)
.Ltmp5:
.Lfunc_end0:
.size _Z3foov, .Lfunc_end0-_Z3foov
.cfi_endproc
```

WORKS WITH LLVM IR.

```
; Function Attrs: uwtable
define void @_Z3foov() #0 {
entry:
    %call = call i32 @printf(i8* getelementptr inbounds ([14 x i8], [14 x i8]* @.str, i32 0, i32 0))
    ret void
}
```

```
attributes #0 = { uwtable "disable-tail-calls"="false" "function-instrument"="xray-always" "less-precise-fpmad"="false" "no-frame-pointer-elim"="true" "no-frame-pointer-elim-non-leaf" "no-infs-fp-math"="false" "no-jump-tables"="false" "no-nans-fp-math"="false" "stack-protector-buffer-size"="8" "target-cpu"="x86-64" "target-features"="+fxsr,+mmx,+sse,+sse2,+x87" "unsafe-fp-math"="false" "use-soft-float"="false" }
```


"function-instrument" is the keyword attribute coming from source-level attributes in C/C++, may be either **"xray-always"** or **"xray-never"**.

automatic function size heuristic comes in with **"xray-instruction-threshold"="NNN"**.

automatic function size heuristic comes in with **"xray-instruction-threshold"="NNN"**.

THE INSTRUMENTATION MAP

- ▶ At runtime we need to know where the sleds are, and what kind of sleds they are in the binary.
- ▶ We keep track of those in the object files.
- ▶ We also know whether to always or sometimes instrument them.
- ▶ We keep some bytes for future use later, and to align the entries better.
- ▶ They are concatenated together by the linker.



THE INSTRUMENTATION MAP

```
.p2align 4, 0x90
.quad .Lxray_synthetic_0
.section xray_instr_map,"a",@progbits
.Lxray_synthetic_0:
.quad .Lxray_sled_0
.quad _Z3foov
.byte 0
.byte 1
.zero 14
.quad .Lxray_sled_1
.quad _Z3foov
.byte 1
.byte 1
.zero 14
.text
```

Our XRay instrumentation map section.

Our entry sled for function "void foo()", which should always be instrumented.

Our exit sled for function "void foo()", which should always be instrumented.

**RUNTIME
LIBRARY.**

PATCHING/UNPATCHING

- ▶ Read through the instrumentation map.
- ▶ Compute the function id (function appearance order) and get the offset of the `__xray_Function{Enter,Exit}` from the sled.
- ▶ For entry sleds, make them look like:
"mov <function id>, %r10d; call __xray_FunctionEnter"
- ▶ For exit sleds, make them look like:
"mov <function id>, %r10d; jmp __xray_FunctionExit"

BEFORE PATCHING

```
(gdb) disassemble foo
```

```
Dump of assembler code for function foo():
```

```
0x000000000000415320 <+0>:      jmp      0x41532b <foo()+11>
0x000000000000415322 <+2>:      nopw     0x200(%rax,%rax,1)
0x00000000000041532b <+11>:     push    %rbp
0x00000000000041532c <+12>:     mov     %rsp,%rbp
0x00000000000041532f <+15>:     sub     $0x10,%rsp
0x000000000000415333 <+19>:     movabs  $0x41ab2c,%rdi
0x00000000000041533d <+29>:     mov     $0x0,%al
0x00000000000041533f <+31>:     callq   0x401a50 <printf@plt>
0x000000000000415344 <+36>:     mov     %eax,-0x4(%rbp)
0x000000000000415347 <+39>:     add     $0x10,%rsp
0x00000000000041534b <+43>:     pop     %rbp
0x00000000000041534c <+44>:     retq
0x00000000000041534d <+45>:     nopw     %cs:0x200(%rax,%rax,1)
```

```
End of assembler dump.
```


AFTER PATCHING

(gdb) disassemble foo


Dump of assembler code for function foo():

```
0x0000000000415320 <+0>:      mov     $0x1,%r10d
0x0000000000415326 <+6>:      callq  0x4145e0 <__xray_FunctionEntry>
0x000000000041532b <+11>:     push   %rbp
0x000000000041532c <+12>:     mov     %rsp,%rbp
0x000000000041532f <+15>:     sub     $0x10,%rsp
0x0000000000415333 <+19>:     movabs  $0x41ab2c,%rdi
0x000000000041533d <+29>:     mov     $0x0,%al
0x000000000041533f <+31>:     callq  0x401a50 <printf@plt>
0x0000000000415344 <+36>:     mov     %eax,-0x4(%rbp)
0x0000000000415347 <+39>:     add     $0x10,%rsp
0x000000000041534b <+43>:     pop     %rbp
0x000000000041534c <+44>:     mov     $0x1,%r10d
0x0000000000415352 <+50>:     jmpq    0x4146d0 <__xray_FunctionExit>
```

End of assembler dump.

LOGGING FUNCTION ENTRY/EXITS

- ▶ The provided logging function gets called from the trampolines.
- ▶ The default implementation logs the timestamp (TSC), thread id, the cpu id, and whether it was a function entry or an exit.
- ▶ We can then reconstruct the function call stack offline based on the entry/exit events in the log.



CUSTOM FUNCTIONALITY

- ▶ We can install any handler with `__xray_set_handler(...)`.
 - ▶ The handler only needs to take two arguments: function id (`int32_t`) and entry type (an enum type, i.e. `int`).
 - ▶ The handler has to be thread-safe.
 - ▶ The handler can pretty much do what it wants. :)
- ▶ Install other handlers with `__xray_set_handler_arg1(...)` for capturing first argument of functions attributed to capture the first argument (useful for "this" pointer)
- ▶ Support for custom event logging with `__xray_set_customevent_handler(...)` for capturing custom events provided by `__xray_customevent(...)` clang built-in.

ANALYSIS TOOLS.

STATISTICS AND RECONSTRUCTION

- ▶ Analysis tools are currently part of the LLVM tools distribution. The tool is called 'llvm-xray' which includes:
 - ▶ Function call accounting statistics.
 - ▶ Function call graph with latency distributions.
 - ▶ Function call stacks with latency sums and counts.

PRACTICAL MATTERS

HOW TO USE XRAY TODAY

USING XRAY IN LLVM

- ▶ Top of Trunk Clang + LLVM + compiler-rt
- ▶ Top of Trunk LLVM, by adding attributes to LLVM IR
- ▶ Available for Linux running on x86_64, ARM7 and ARM8 (no thumb), AArch64, PowerPC 64 (little endian), MIPS

DOCUMENTATION FOR XRAY

- ▶ XRay [white paper](#)
- ▶ High level documentation (llvm.org/docs/XRay.html)
- ▶ Debugging Example (llvm.org/docs/XRayExample.html)

AN EXAMPLE IN C++

CODE AND TOOLS

HELLO, XRAY! (CODE)

```
#include <iostream>
```

```
[[clang::xray_always_instrument]] void foo() {  
    std::cout << "Hello, XRay!" << std::endl;  
}
```

```
[[clang::xray_always_instrument, clang::xray_log_args(1)]]  
void bar(int i) {  
    std::cout << "Captured: " << i << std::endl;  
}
```

```
[[clang::xray_always_instrument]] int main(int argc, char* argv[]) {  
    foo();  
    bar(argc);  
}
```

EXAMPLE

HELLO, XRAY! (BUILD)

```
clang++ -fxray-instrument -std=c++11 \  
        hello_xray.cpp -o hello_xray
```

EXAMPLE

HELLO, XRAY! (RUN)

```
$ ./hello_xray  
Hello, XRay!  
Captured: 1
```

```
$ XRAY_OPTIONS="patch_premain=true" ./hello_xray  
==NNNNN==XRay: Log file in 'xray-log.hello_xray.xxxx'  
Hello, XRay!  
Captured: 1
```

HELLO, XRAY! (TOOLS)

```
$ llvm-xray stack -instr_map=./hello_xray xray-log.hello_xray.*
Unique Stacks: 2
Top 10 Stacks by leaf sum:
```

Sum: 124448

lvl	function	count	sum
#0	main	1	176552
#1	foo()	1	124448

Sum: 43016

lvl	function	count	sum
#0	main	1	176552
#1	bar(int)	1	43016

Top 10 Stacks by leaf count:

Count: 1

lvl	function	count	sum
#0	main	1	176552
#1	foo()	1	124448

Count: 1

lvl	function	count	sum
#0	main	1	176552
#1	bar(int)	1	43016

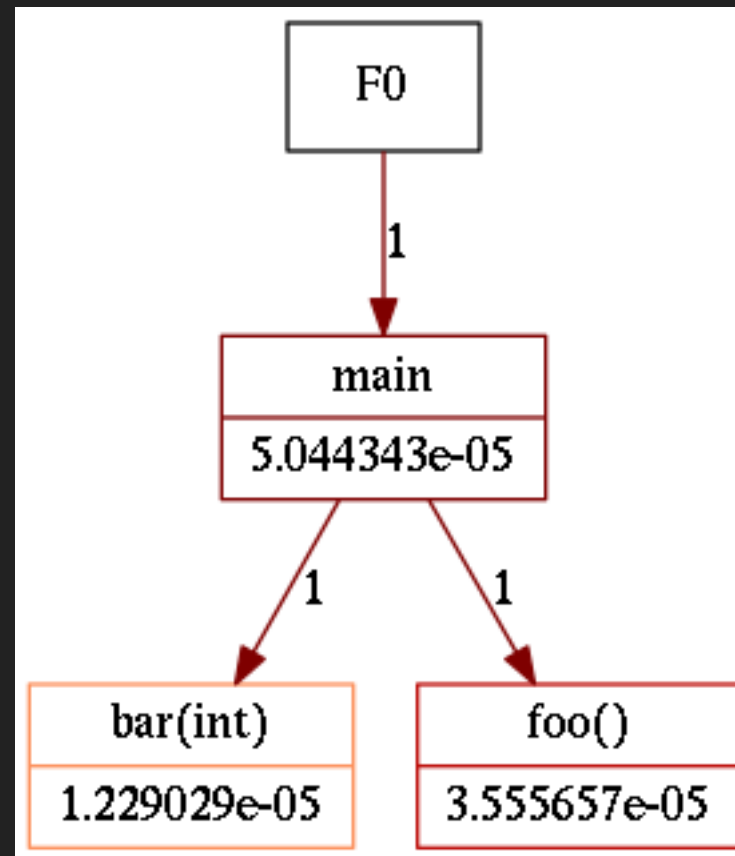
HELLO, XRAY! (TOOLS)

```
$ llvm-xray account -instr_map=./hello_xray xray-log.hello_xray.*
Functions with latencies: 3
funcid      count [      min,      med,      90p,      99p,      max]      sum function
    1         1 [ 0.000036, 0.000036, 0.000036, 0.000036, 0.000036] 0.000036 <invalid>:0:0: foo()
    2         1 [ 0.000012, 0.000012, 0.000012, 0.000012, 0.000012] 0.000012 <invalid>:0:0: bar(int)
    3         1 [ 0.000050, 0.000050, 0.000050, 0.000050, 0.000050] 0.000050 <invalid>:0:0: main
```


EXAMPLE

HELLO, XRAY! (TOOLS)

```
$ llvm-xray graph -instr_map=./hello_xray xray-log.hello_xray.* -color-edges=count \
  -edge-label=count -color-vertices=sum -vertex-label=sum | dot -png -x > /tmp/hello_xray.png
```



EXAMPLE

HELLO, XRAY! (TOOLS)

```
$ llvm-xray stack -instr_map=./hello_world_fdr \  
  xray-log.hello_world_fdr.s3c4ft --keep-going --stack-format=flame \  
  -all-stacks 2>/dev/null | ~/Source/FlameGraph/flamegraph.pl \  
  > /tmp/hello_xray_flame.svg
```

FlameGraph tool from
<https://github.com/brendangregg/FlameGraph>

DEMO.

FOR THE DEVS.

XRAY IMPLEMENTATION DETAILS

XRAY IN LLVM

LLVM

- ▶ XRay Instrumentation Pass

Takes MachineFunctions and inserts pseudo instructions marking where the start of the function is, where the exits are, and any custom event points.

- ▶ CodeGen implementations for targets

Lowering to per-target specific assembly, ELF/MachO sections for the instrumentation map.

- ▶ Analysis tools and libraries

Implementation of the llvm-xray tool with various sub-commands (extract, convert, stacks, account, graph, graph-diff).

Trace reader library for reading XRay traces.

InstrumentationMap class for working with the instrumentation map in XRay-instrumented binaries.

XRAY IN CLANG

▶ Driver Options

Support for `-fxray-instrument`, `-fxray-instruction-threshold`, `-fxray-always-instrument=...`, and `-fxray-never-instrument=...` -- also adding the compiler runtime library when linking binaries.

▶ Builtin Support

Support for `__xray_customevent(...)`.

▶ Source-level Attributes (C/C++)

We use C++ attributes, `[[clang::xray_always_instrument]]`, `[[clang::xray_never_instrument]]`, and `[[clang::xray_log_args(N)]]` and `__attribute__((...))` analogues of these.

XRAY IN COMPILER-RT

- ▶ XRay APIs and Implementation

Control APIs for installing handlers, turning instrumentation on/off, etc.

- ▶ Logging Implementations

Basic (naive) mode and Flight Data Recorder mode.

SUPPORT MATRIX

feature	x86_64	ppc64le	arm	aarch64	mips
no-arg logging	✓	✓	✓	✓	✓
1-arg logging	✓	✓	✓	✓	✓
basic/naive mode	✓	✓	✓	✓	✓
custom events	✓				
flight data recorder mode	✓				

**WHAT TO EXPECT
NEXT.**

- ▶ More Analysis Tools
- ▶ Support for more function exit types (exceptional exits, interrupts (maybe), etc.)
- ▶ More architectures and platforms/OSes

THANK YOU!

QUESTIONS?

Dean Berris ([@deanberris](#)), XRay Team (google-xray@googlegroups.com)