

Portable yet thin OS abstractions

Klemens Morgenstern


18.10.2018

- Electrical engineer by training
- C++ developer by passion
- Author of boost.process
- Independent contractor
- Open for consulting

Goal of C++ system libraries

- Portable C++ Applications
- Low/Zero overhead
- Thin rather than thick
- Libraries and not frameworks

Portable C++ Application



- Application code compiles against different OS
- OS specific extensions are possible

'In general, C++ implementations obey the zero-overhead principle: What you don't use, you don't pay for. And further: What you do use, you couldn't hand code any better.'

Bjarne Stroustrup

- OS stream: byte sequence with `write()` and `read()`
- C++ stream: byte sequence with formatting, locale, `iosmanip` to an OS stream or memory
- C++ streams are a replacement of `xprintf` & `xscanf`
- Thick as a standard I/O system library
- Do not give access to the handles (OS layer)

- Only pay for what you use
- Do not enforce a design
- Do not conflict with other libraries
- Do not hide the OS

Discussed libraries

boost.thread	<thread>	C++11
boost.asio	Networking	TS
boost.process		
boost.filesystem	<filesystem>	C++17

Operating Systems

Windows

Posix Linux, Free BSD, Solaris, OSX

`std::thread`

- Based on design of `boost::thread`
- Standardized in C++11
- An RAII wrapper around OS threads

C++ thread expected code

```
1  class thread {
2      native_handle _handle;
3  public:
4      using native_handle = undefined;
5
6      template<typename Func>
7      thread(Func && func);
8
9      thread(thread &&);
10     thread& operator=(thread &&);
11
12     void join();
13     void terminate();
14     void detach();
15
16     ~thread();
17 };
```

C++ thread expected code (Win32)

```
1  #if defined(_WIN32)
2  template<typename Func>
3  void thread::thread(Func&& func)
4      : _handle(CreateThread (0,0, +[](Func && f){f();}, &func, 0)) {}
5
6  void thread::join() {
7      WaitForSingleObject(_handle, INFINITE);
8      _handle = nullptr;
9  }
10
11 void thread::terminate() {
12     TerminateThread(_handle, 0);
13     _handle = nullptr;
14 }
15
16 void thread::detach() { _handle = nullptr; }
17
18 thread::~~thread() {
19     if (_handle != nullptr)
20         terminate();
21 }
22 #else
```

C++ thread expected code (posix)

```
23  template<typename Func>
24  void thread::thread(Func&& func)
25  {
26      pthread_create(&_amp;handle, nullptr, +[](Func *f){(*f)();}, &func);
27  }
28  void thread::join() {
29      pthread_join(_amp;handle, nullptr);
30      _amp;handle = 0;
31  }
32
33  void thread::terminate() {
34      pthread_cancel(_amp;handle);
35      _amp;handle = 0;
36  }
37
38  void thread::detach() { _amp;handle = 0; }
39
40  thread::~~thread() {
41      if (_amp;handle != 0)
42          terminate();
43  }
44  #endif
```

Boost::thread & std::thread

```
1  class thread {
2      native_handle _handle;
3  public:
4      template<typename Func>
5      thread(Func && func);
6
7      thread(thread &&);
8      thread& operator=(thread &&);
9      // void terminate();
10     void join();
11     void detach();
12     ~thread() {
13         if (_handle)
14             std::terminate();
15     }
16 };
```

- No terminate function?
- std::terminate() in destructor?

Boost::thread & std::thread - rationale

```
1  mutex mtx;
2
3  thread thr{
4      [&]{
5          lock_guard<mutex> lock{mtx};
6          while(true)
7              this_thread::sleep_for(chrono::seconds(1));
8      };
9
10 this_thread::sleep_for(chrono::seconds(5));
11
12 thr.terminate();
13 lock_guard<mutex> lock(mtx);
14 std::cerr << "Is_this_line_reached?" << std::endl;
```

- On windows, line 14 is never reached (no stack unwinding)
- On posix (pthreads), stack unwinding releases the mutex
- Different semantics between Windows / Posix

boost::thread & std::thread - why std::terminate?

- Destructor can't throw
- RAII requires a terminate of the thread

Terminating a thread anyhow

```
1 struct my_thread : std::thread {
2     void terminate() {
3         #if defined(_WIN32)
4             TerminateThread(native_handle(), 0);
5         #else
6             pthread_cancel(native_handle());
7         #endif
8     }
9     ~thread() {
10         if (joinable()) {
11             terminate();
12             join();
13         }
14     }
15 }
```

boost::thread::interrupt extension

```
1  using namespace boost;
2  thread thr{
3      [&]{
4          try {
5              unique_lock<mutex> lock(mtx);
6              condition_variable cond;
7              cond.wait(lock);
8          }
9          catch (boost::thread_interrupted&) {}
10     };
11
12     thr.interrupt();
13     thr.join();
```

- Consistent semantics
- Not implemented in the OS
- Overhead when not used

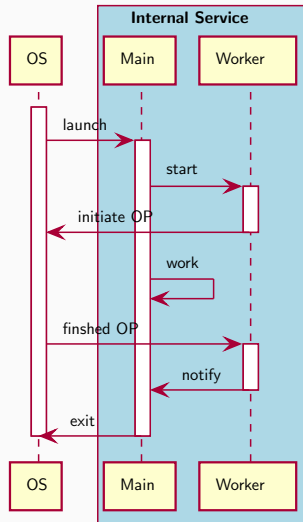
boost::asio

- Basis for the Networking TS
- Provides facilities for sync and async I/O
- TCP, UDP, SerialPort

Asynchronous I/O with threads

```
1 auto res =  
2     std::async(std::launch::async ,  
3               os::sync_op );  
4  
5 while(res.wait_for(1_ms)  
6       == future_status::timeout)  
7     do_work();  
8  
9 clog << res.get() << endl;
```

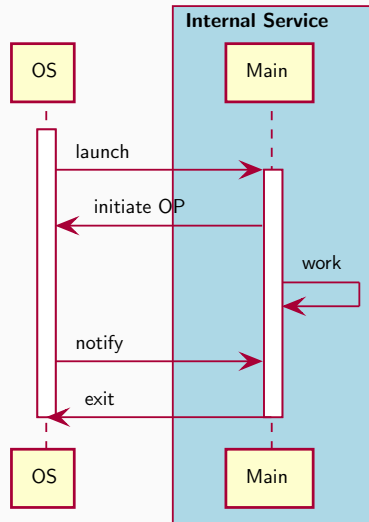
- Threads cause overhead
- One Thread per Stream



Asynchronous I/O with Syscalls

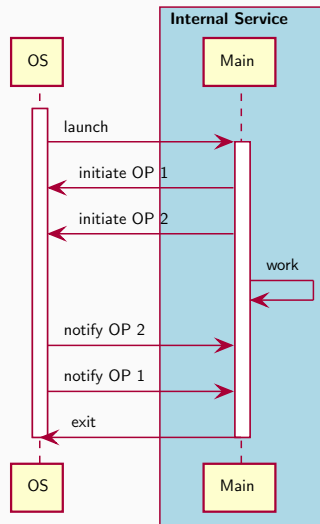
```
1 auto res = os::async_op();  
2  
3 while(res.wait_for(1_ms)  
4       == os::timeout)  
5     do_work();  
6  
7 clog << res.get() << endl;
```

- No overhead
- One Thread for all OPs
- Better design for more OPs



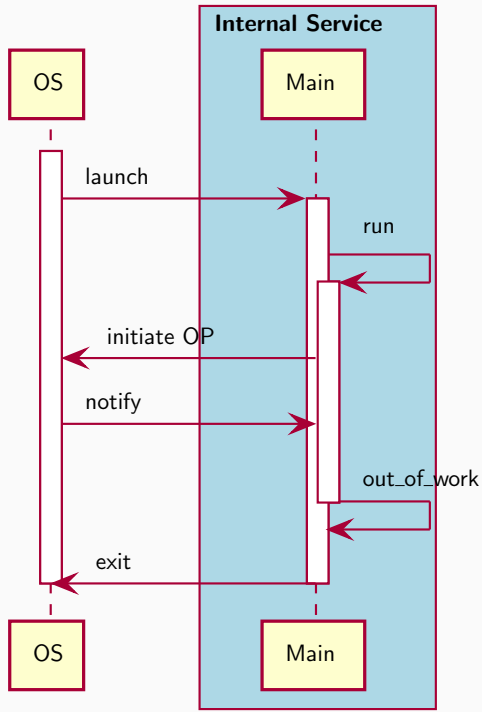
Asynchronous I/O with Syscalls

```
1 auto res1 = os::async_op1();
2 auto res2 = os::async_op2();
3
4 while((res1.wait_for(1_ms)
5       == os::timeout) ||
6       (res2.wait_for(1_ms)
7       == os::timeout))
8     do_work();
9
10 clog << res1.get()
11    << res2.get() << endl;
```



asio example

```
1 using resolver = asio::tcp::resolver;  
2 asio::io_context ioc;  
3  
4 resolver resolv{ioc};  
5 resolver::query q{  
6     "pacificplusplus.com",  
7     443};  
8  
9 auto fut = resolv.async_resolve(  
10     q,  
11     asio::use_future);  
12  
13 ioc.run();  
14  
15 auto res = fut.get();
```



System Functionality

- Descriptors/Handles
 - Posix: file descriptor (`int`)
 - Windows: stream handle (`void*`)
 - Posix: read/write
 - Windows: `ReadFile(Ex)/WriteFile(Ex)`
 - Socket/FileStream/Pipe/SerialPort
- Waitables / Events
 - Posix: signal (`void(*) (int)`)
 - Windows: event handle (`void*`)

Usage examples

- Child Process Changes
- Interrupt request
- Timers

```
1 static std::function<void()> handler;  
2  
3 using signal_type = void(*)(int);  
4 signal_type old_sig = nullptr;  
5  
6 handler = [&]{cout << "signal_received" << endl; old_sig();};  
7 old_sig = signal(SIGTERM, +[](int){handler();});
```

- Similar to interrupts, i.e. stop execution
- Should just notify, not do work
- Signal handler blocks signal
- Do not scale well

Examples

- Process Handles
- Thread Handles
- Events
- Querys, e.g. hostname lookup

Functions

- `WaitForSingleObject(Ex)`
- `WaitForMultipleObjects(Ex)`
- `RegisterWaitForSingleObject`
- `UnregisterWait(Ex)`

Event & RegisterWaitForSingleObject

```
1  auto handler = []{ std::cerr << "Event_triggered" << std::endl;};
2
3  auto event = CreateEventA(nullptr, false, false, "Pacific++");
4  HANDLE wait_handle;
5  //May spawn a thread pool
6  RegisterWaitForSingleObject(
7      &wait_handle,
8      event,
9      +[](void *p, unsigned char)
10     {
11         auto h = static_cast<decltype(handler)*>(p);
12         (*h)();
13     },
14     &handler,
15     INFINITE,
16     WT_EXECUTEONLYONCE);
17
18  SetEvent(event);
```

Overlapped & Events

```
1 HANDLE file_handle = CreateFile("pacifi.c++",
2                                 GENERIC_READ|GENERIC_WRITE,
3                                 0, NULL, OPEN_EXISTING,
4                                 FILE_ATTRIBUTE_NORMAL|
5                                 FILE_FLAG_OVERLAPPED, NULL);
6
7 OVERLAPPED overlapped{};
8
9 //Sync version GetOverlappedResult --> no thread-pool
10 overlapped.hEvent = CreateEvent(NULL, TRUE, FALSE, NULL);
11 RegisterWaitForSingleObject(overlapped/*,...*/); //previous slide
12
13 auto res =
14     WriteFile(file_handle,
15              "void_foo();", 11,
16              NULL, &ovlp);
17
18 if (!res && GetLastError() != ERROR_IO_PENDING) { /*errored*/ }
```

Overlapped & I/O Completion ports

```
1  HANDLE file_handle = CreateFile(/*...*/);
2  OVERLAPPED overlapped{};
3
4  auto io_port = CreateIoCompletionPort(
5      file_handle, NULL, //ExistingCompletionPort
6      NULL, 0);
7
8  auto status = WriteFile(/*...*/ &ovlp);
9
10 OVERLAPPED * ovl_p;
11 DWORD bytes_transferred = 0;
12 ULONG_PTR completion_key = 0;
13
14 if (GetQueuedCompletionStatus(
15     io_port, &bytes_transferred,
16     &completion_key, &ovl_p, INFINITE)) {
17 }
```

Overlapped & I/O Completion ports

```
1  HANDLE file_handle = CreateFile(/*...*/);
2  struct my_overlapped : OVERLAPPED { int foo = 42;};
3  my_overlapped overlapped{};
4
5  auto io_port = CreateIoCompletionPort(
6      fh, NULL, //ExistingCompletionPort
7      NULL, 0);
8
9  auto status = WriteFile(/*...*/ &ovlp);
10
11  OVERLAPPED * ovl_p;
12  DWORD bytes_transferred = 0;
13  ULONG_PTR completion_key = 0;
14
15  if (GetQueuedCompletionStatus(
16      io_port, &bytes_transferred,
17      &completion_key, &ovl_p, INFINITE)) {
18      static_cast<my_overlapped*>(ovl_p)->foo = 12;
19  }
```

Summary Windows I/O

- Callbacks for events
- Callback based async I/O per Operation
- Callback may spawn a thread-pool
- Completion Ports for querying and waiting for multiple OPs

Possible on Windows (using the OS thread-pool)

```
1 auto fut = fd.async_write("foo");  
2 fut.then([fd&]{fd.async_write("bar");});
```

Asynchronous I/O on Posix

- Posix aio.h library
 - Callback based
 - Only supports File I/O
 - Used in boost.afio
- SIGIO
 - Not posix, but some posix implementations
 - Only supports a subset of streams (e.g. sockets, but not pipes)
 - Brings all problems of signals
- **A posix implementation cannot be callback based**

Not possible on Posix (without background threads)

```
1 auto fut = fd.async_write("foo");  
2 fut.then([fd&]{fd.async_write("bar");});
```

select/poll Posix

```
1  fcntl(STDIN_FILENO, F_SETFL, O_NONBLOCK);
2  fcntl(STDOUT_FILENO, F_SETFL, O_NONBLOCK);
3  fcntl(STDERR_FILENO, F_SETFL, O_NONBLOCK);
4
5  struct pollfd fds[3] = {
6      {.fd = STDIN_FILENO, .events = POLLIN},
7      {.fd = STDOUT_FILENO, .events = POLLOUT},
8      {.fd = STDERR_FILENO, .events = POLLOUT}
9  };
10
11  write(STDOUT_FILENO, "Some_data", 9);
12  write(STDERR_FILENO, "More_data", 9);
13
14  auto fd_cnt = ::poll(fds, 3, -1);
```

- select does essentially the same
- POLLOUT triggers when handle is available for a write
- POLLIN triggers when data is available for read

select/poll alternatives

select/poll are slow, because they check all descriptors

OS	Mechanism	Used by asio?
General Posix	poll/select	select_reactor
Linux	epoll	epoll_reactor
BSD	kqueue	kqueue_reactor
Solaris	/dev/poll	dev_poll_reactor
AIX	pollset	

Common functionality?

	Windows	Posix
Callbacks as notifications	Yes	Yes
Callbacks for work	Yes	No
Callback per stream	Yes	No
Poll multiple streams	Yes	Yes
Wait for multiple streams	Yes	Yes

- Combining all stream-handles / -descriptors
- poll or IoCompletionPort
- Polling: asio::io_context::poll
- Waiting: asio::io_context::run
- Notification/Callback-handling?
 - PostQueuedCompletionStatus on windows
 - Event File-Descriptor on posix (self-pipe or eventfd)
- Queue work manually with io_context::post
- **Introduced a new concept, the io_context**
- Little overhead regarding polling
- Requires queueing of operations

asio::io_context

```
1 namespace bp = boost::process;
2
3 asio::io_context ioc;
4 bp::async_pipe apipe{ioc};
5
6 std::string write_buf = "Pacific++";
7 std::string read_buf;
8
9 std::future<std::size_t> read_fut;
10 auto write_handler =
11     [](auto ec, auto sz) {
12         read_buf.resize(sz);
13         read_fut = asio::async_read(apeipe,
14                                     asio::buffer(read_buf),
15                                     asio::use_future);
16     };
17 asio::async_write(apeipe, asio::buffer(write_buf), write_handler);
18
19 ioc.run(); //Start poll/GetQueuedCompletionStatus
```

boost::process

- Provides management for processes
- Pipes & Environment

Starting a process

Possible parameters for starting a process

- Command & Args
- Redirection STDIN, STDOUT, STDERR
- Working directory
- Environment
- Platform extensions
- User extensions (at startup)

Variadic interface

```
1 namespace bp = boost::process;  
2 child c1{bp::exe="./my_app", bp::args={"foo", "bar"},  
3     bp::start_dir="./working_dir"};  
4  
5 bp::pipe p_stdout;  
6 child c2{bp::exe="./other_app", bp::std_out=p_stdout};
```

Launching a process on Windows

```
1 namespace bp = boost::process;  
2 child c1{bp::exe="./my_app", bp::args={"foo", "bar"},  
3         bp::start_dir="./working_dir"};
```

```
1 STARTUP_INFO startup_info {...};  
2 PROCESS_INFORMATION proc_info;  
3 CreateProcessW(  
4     L"my_app.exe", L"foo\0bar\0",  
5     nullptr, nullptr, //process & thread attributes  
6     FALSE, //inherit handles  
7     0, //creation flags  
8     nullptr, //environment  
9     L"./working_dir",  
10    &startup_info,  
11    &proc_info);
```

Launching a process on Posix

```
1 namespace bp = boost::process;  
2 child c1{bp::exe="./my_app", bp::args={"foo", "bar"},  
3       bp::start_dir="./working_dir"};
```

```
1 auto pid = ::fork();  
2  
3 if (pid == -1) { /* error */ }  
4 else if (pid == 0) //child process  
5 {  
6     chdir("./working_dir");  
7     const char * args[] = {"foo", "bar", nullptr};  
8     execv("./my_app", args);  
9 }
```

Launching a process on Posix with stdout

```
1 bp::pipe p_stdout;  
2 child c2{bp::exe="./other_app", bp::std_out=p_stdout);
```

```
1 bp::pipe p_stdout;  
2  
3 auto pid = ::fork();  
4  
5 if (pid == -1) { /* error */ }  
6 else if (pid == 0) //child process  
7 {  
8     dup2(p_stdout.native_sink(), STDOUT_FILENO);  
9     p_stdout.close();  
10     const char * args[] = {nullptr};  
11     execv("./other_app", args);  
12 }  
13  
14 close(p_stdout.native_sink());
```


Launching a process on Windows with stdout

```
1  bp::pipe p_stdout;  
2  SetHandleInformation(p_stdout.native_sink(),  
3                      HANDLE_FLAG_INHERIT,  
4                      HANDLE_FLAG_INHERIT);  
5  
6  STARTUP_INFO startup_info;  
7  startup_info.hStdOutput = p_stdout.native_sink();  
8  startup_info.dwFlags |= STARTF_USESTDHANDLES;  
9  
10 PROCESS_INFORMATION proc_info;  
11 CreateProcessW(  
12     L"my_app.exe", L"\0\0",  
13     nullptr, nullptr, //process & thread attributes  
14     TRUE, //inherit handles  
15     0, //creation flags  
16     nullptr, //environment  
17     nullptr, //Working dir  
18     &startup_info,  
19     &proc_info);  
20  
21 CloseHandle(p_stdout.native_sink());
```

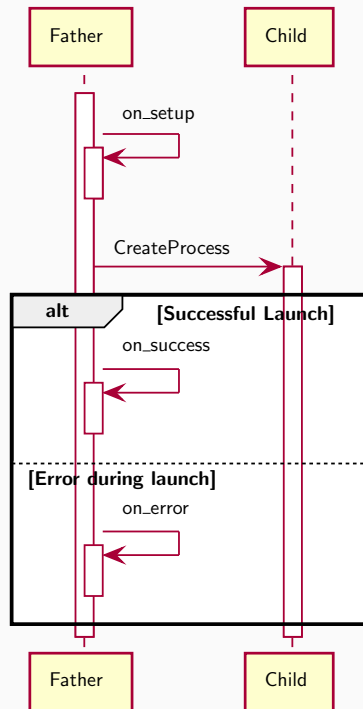
Implementation with Initializers

- Arguments build an initializers tuple
- Executor takes the tuple and executes different steps

```
1  template<typename ... Args> child::child(Args && ... args)
2      : child(make_executor(std::tie(args))()) {}
3
4  template<typename Seq> struct executor {
5      Seq seq;
6      void operator()() {
7          for_each(seq, [this](auto & e){e.on_setup(*this);})
8          auto res = CreateProcess(...);
9
10         if (!res) {
11             auto err = std::get_last_error();
12             for_each(seq, [&](auto & e){e.on_error(*this, err);})
13         }
14         else
15             for_each(seq, [this](auto & e){e.on_success(*this);})
16         return child(handle);
17     }
18 };
```

```
1  struct handler {  
2  
3      template <class Executor>  
4      void on_setup(Executor&) const {}  
5  
6      template <class Executor>  
7      void on_error(Executor&, const std::error_code &) const {}  
8  
9      template <class Executor>  
10     void on_success(Executor&) const {}  
11 };
```


Executor structure windows



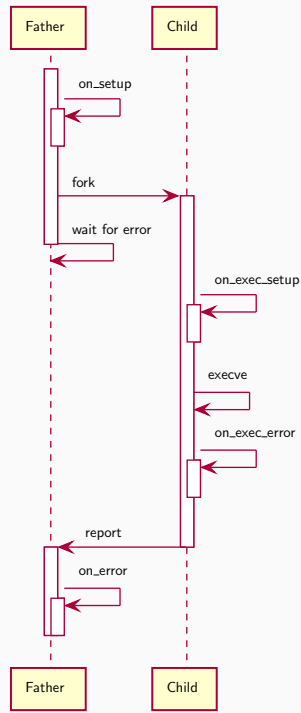

Posix Handler

```
1 struct handler {
2     template <class Executor>
3     void on_setup(Executor&) const {}
4
5     template <class Executor>
6     void on_error(Executor&, const std::error_code &) const {}
7
8     template <class Executor>
9     void on_success(Executor&) const {}
10
11     template<typename Executor>
12     void on_fork_error(Executor &, const std::error_code&) const {}
13
14     template<typename Executor>
15     void on_exec_setup(Executor &) const {}
16
17     template<typename Executor>
18     void on_exec_error(Executor &, const std::error_code&) const {}
19 };
```

Executor structure posix



Executor structure posix



start_dir posix initializer

```
1 namespace boost::process::detail::posix
2 {
3
4 struct start_dir_init : handler
5 {
6     start_dir_init(const std::string &s) : s_(s) {}
7
8     template <class PosixExecutor>
9     void on_exec_setup(PosixExecutor&) const
10     {
11         ::chdir(s_.c_str());
12     }
13
14     const std::string & str() const {return s_;}
15 private:
16     std::string s_;
17 };
18
19 }
```


start_dir windows initializer

```
1 namespace boost::process::detail::windows
2 {
3
4 struct start_dir_init : handler
5 {
6     start_dir_init(const std::string &s) : s_(s) {}
7
8     template <class Executor>
9     void on_setup(Executor& exec) const
10    {
11        exec.work_dir = s_.c_str();
12    }
13
14    const std::string & str() const {return s_;}
15 private:
16     std::string s_;
17 };
18
19 }
```

Extension

```
1 struct handler {
2
3     template <class Executor>
4     void on_setup(Executor&) const {}
5
6     template <class Executor>
7     void on_error(Executor&, const std::error_code &) const {}
8
9     template <class Executor>
10    void on_success(Executor&) const {}
11 };
```

```
1 struct my_extension : boost::process::handler {
2     template <class Executor>
3     void on_setup(Executor &) const {
4         std::cout << "Hello_Pacific++" << std::endl;
5     }
6 };
7
8 bp::child c("foo", my_extension {});
```

async_pipe

```
1 asio::io_context ioc;  
2  
3 bp::pipe p;  
4 bp::async_pipe ap{ioc};
```

- On posix, O_NONBLOCK can be set on any pipe
- Windows only supports OVERLAPPED for named pipes
- async_pipe is always named on windows

Pipe name generation on windows

```
1 std::string make_pipe_name() {  
2     std::string name = R"(\.\pipe\boost_process_auto_pipe_)";  
3     auto pid = GetCurrentProcessId();  
4     static std::atomic_size_t cnt{0};  
5     name += std::to_string(pid);  
6     name += "_";  
7     name += std::to_string(cnt++);  
8     return name;  
9 }
```

- The User should be able to name a pipe
- Posix does not have a named pipe like windows

Named pipe on posix

```
1 std::pair<int, int> make_named_pipe(const std::string &name) {  
2     auto fifo = mkfifo(name.c_str(), 0666 );  
3     int read_fd = open(name.c_str(), O_RDWR);  
4     int write_fd = dup(read_fd);  
5     return {read_fd, write_fd};  
6 }
```

Using a named pipe

```
1 #if defined(BOOST_WINDOWS_API)  
2 std::string pipe_name = R"(\\.\pipe\my_pipe)";  
3 #else  
4 std::string pipe_name = "./my_pipe";  
5 #endif  
6 bp::pipe named_pipe(pipe_name);  
7 bp::async_pipe ap(ioc, named_pipe);
```

- www.github.com/klemens-morgenstern
- klemens.d.morgenstern@gmail.com
- Any questions?