

An Introduction to the Coroutines TS

Toby Allsopp

WhereScape Software Limited

email: toby@mi6.gen.nz

twitter: [@toby_allsopp](https://twitter.com/toby_allsopp)

github: [toby-allsopp](https://github.com/toby-allsopp)

Introduction

Structure

- Context
- Basics
- Motivating examples
- Under the covers

What are we talking about?

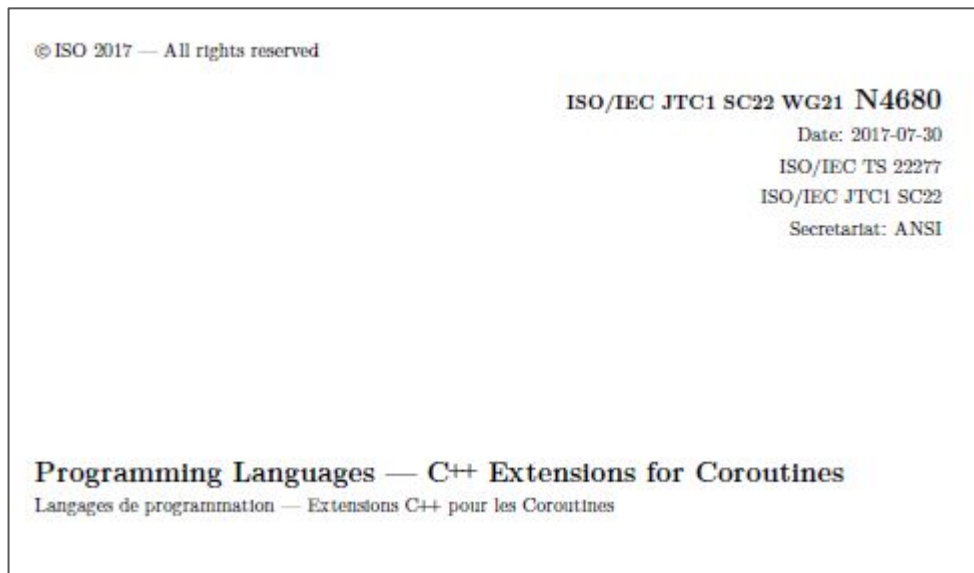
“Programming Languages — C++ Extensions for Coroutines”

aka N4680

a Technical Specification

henceforth: “the TS”

wg21.link/n4680



What's *that* talking about?



Coroutines, obviously!

An extension to C++

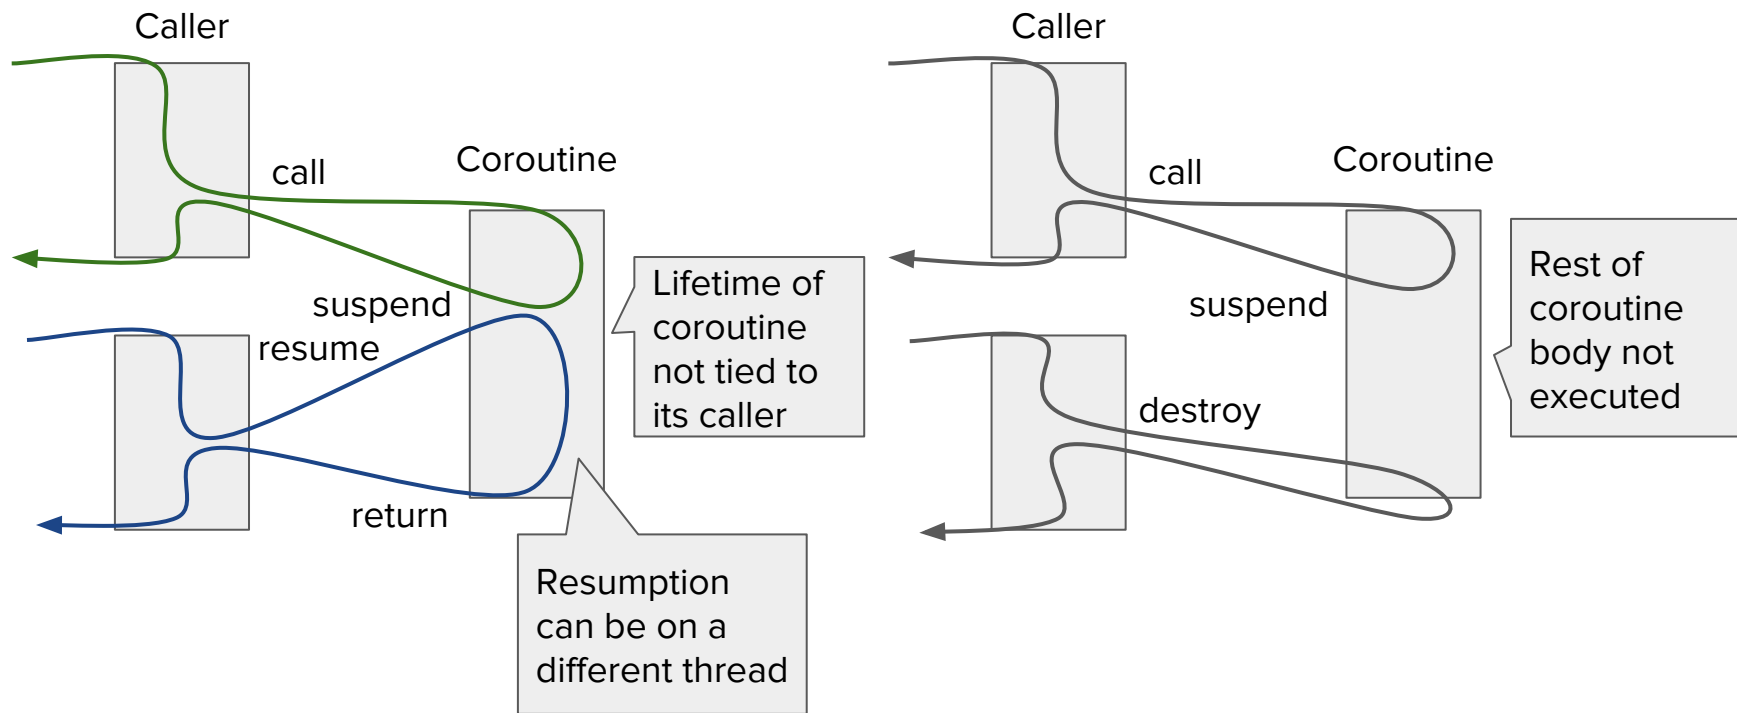
Three new keywords and some new customization points

That enable some new (to C++) kinds of control flow

A coroutine is a generalization of a function



A coroutine is a generalization of a function



Motivating use-cases

Motivating use-cases

| | Synchronous | Asynchronous |
|-------------------------------|--------------------|---------------------|
| Single return value | Normal function | Task |
| Multiple return values | Generator | Async generator |

Generators

A generator:

- computes a sequence of values
- one at a time
- on demand

Generators

Caller

```
bool caller() {  
    for (int p : primes()) {  
        use(p);  
        if (p > 100) return true;  
        if (finished(p)) break;  
    }  
    // many lines...  
    return false;  
}
```

Both caller and
generator want
to be in the
driver's seat

Generator

```
vector<int> primes() {  
    vector<int> ps;  
    prime_tester tester;  
    for (int i = 2;; ++i) {  
        if (tester.is_prime(i))  
            ps.push_back(i);  
    }  
    return ps; // unreachable  
}
```

Clearly, this
won't work!

Take 1 — callback

Caller

```
bool caller() {  
    bool return_early = false;  
    primes([&return_early](int p) {  
        use(p);  
        if (p > 100) {  
            return_early = true;  
            return true; // finished  
        }  
        return finished(p);  
    });  
    if (return_early) return true;  
    // many lines...  
    return false;  
}
```

The generator
is in the
driver's seat

The caller is
messed up

Generator

```
template <class F>  
void primes(F&& callback) {  
    prime_tester tester;  
    for (int i = 2;; ++i)  
        if (tester.is_prime(i))  
            if (callback(i)) return;  
}
```

Take 2 — iterator

Caller

```
bool caller() {  
    for (int p : primes()) {  
        use(p);  
        if (p > 100) return true;  
        if (finished(p)) break;  
    }  
    // many lines...  
    return false;  
}
```

The caller is in
the driver's
seat

The generator
is inside out

Generator

```
#include <range/v3/all.hpp>  
using namespace ranges;
```

```
auto primes() {  
    class range : public view_facade<range> {  
        friend range_access;  
        prime_tester tester_;  
        struct cursor {  
            range* range_;  
            int i_ = 2;  
  
            void next() {  
                do {  
                    ++i_;  
                } while (!range_->tester_.is_prime(i_));  
            }  
            bool done() const { return false; }  
            int get() const { return i_; }  
        };  
  
        cursor begin_cursor() {  
            return cursor{this, 2};  
        }  
    };  
    return range{};  
}
```

Take 3 — coroutine

Caller

```
bool caller() {  
    for (int p : primes()) {  
        use(p);  
        if (p > 100) return true;  
        if (finished(p)) break;  
    }  
    // many lines...  
    return false;  
}
```

Both caller and
generator are
in the driver's
seat

Generator

```
generator<int> primes() {  
    prime_tester tester;  
    for (int i = 2;; ++i) {  
        if (tester.is_prime(i))  
            co_yield i;  
    }  
}
```

`generator<T>` is a
class that uses a
coroutine to compute
values

Async

An asynchronous function:

- returns promptly to its caller (does not block)
- returns an object representing the outstanding work to be done
- provides a way to get the eventual result

(Not) async

- Imagine some API for connecting to and reading from a socket
- It's pretty easy to use if it's **blocking**

```
class Socket {  
    public:  
        template <size_t N>  
        int read(array<char, N>& buf);  
};
```

```
Socket connect(string url);
```

```
string download(string url) {  
    auto s = connect(url);  
    string result;  
    array<char, 1024> buf;  
    while (int n = s.read(buf)) {  
        result.append(buf.data(), n);  
    }  
    return result;  
}
```


(Still not) async

Now imagine the API is non-blocking

```
class Socket {  
    public:  
        template <size_t N>  
        auto read(array<char, N>& buf)  
            -> Task<int>;  
};
```

```
Task<Socket> connect(string url);
```

`Task` is a fictional class template representing a value that may not be available yet

```
string download(string url) {  
    auto s = connect(url).get();  
    string result;  
    array<char, 1024> buf;  
    while (int n = s.read(buf).get()) {  
        result.append(buf.data(), n);  
    }  
    return result;  
}
```

`Task::get()`
blocks until the
value is available

Async (with .then)

```
Task<string> download(string url) {  
    return connect(move(url)).then([](auto socket) {  
        auto state = make_shared<State>(move(socket));  
        return while_task([state]() {  
            return state->socket.read(state->buffer).then([state](int n) {  
                if (n == 0) return make_ready_task(false);  
                state->result.append(state->buffer.data(), n);  
                return make_ready_task(true);  
            });  
        })  
        .then([state]() { return make_ready_task(state->result); });  
    });  
}
```

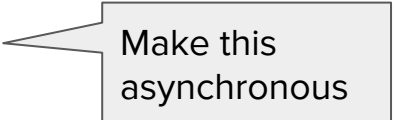
Callbacks turn
the logic inside
out

`Task::then()` returns a new
task that runs a *continuation*
when the value is available

```
struct State {  
    explicit State(Socket socket)  
        : socket(move(socket)) {}  
    Socket socket;  
    array<char, 1024> buffer;  
    string result;  
};
```

Remember what we were trying to do?

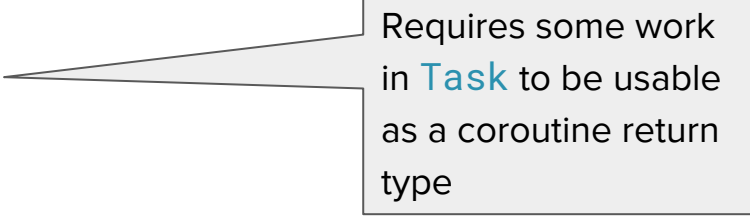
```
    string download(string url) {  
    auto s =          connect(url);  
    string result;  
    array<char, 1024> buf;  
    while (int n =          s.read(buf)) {  
        result.append(buf.data(), n);  
    }  
    return result;  
}
```



Make this
asynchronous

It's easy with coroutines!

```
Task<string> download(string url) {  
    auto s = co_await connect(url);  
    string result;  
    array<char, 1024> buf;  
    while (int n = co_await s.read(buf)) {  
        result.append(buf.data(), n);  
    }  
    co_return result;  
}
```



Requires some work
in `Task` to be usable
as a coroutine return
type

Async generators

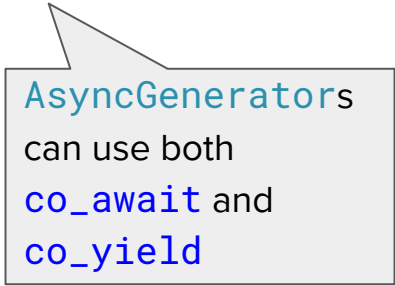
An asynchronous generator:

- computes a sequence of values
- one at a time
- on demand
- asynchronously

Async primes

```
Task<bool> caller()
{
    for co_await (int p : primes)
    {
        use(p);
        if (p > 100) co_return true;
        if (finished(p)) break;
    }
    // many lines...
    co_return false;
}
```

```
AsyncGenerator<int> primes()
{
    async_prime_tester tester;
    for (int i = 2;; ++i)
    {
        if (co_await
            tester.is_prime(i))
            co_yield i;
    }
}
```




AsyncGenerators
can use both
`co_await` and
`co_yield`

Coroutine transformations and customization points


Very simple coroutine

```
void caller() {  
    auto lazy_result = calculate();  
}
```



```
lazy<int> calculate() {  
    int result = hard_work();  
    co_return result;  
}
```

```
int caller() {  
    auto lazy_result = calculate();  
    return lazy_result.get();  
}
```



```
lazy<int> calculate() {  
    int result = hard_work();  
    co_return result;  
}
```


Coroutine body transformation

```
lazy<int> calculate()
```

```
{
```

the promise is
actually stored
in the coroutine
state

```
int result = hard_work();  
co_return result;
```

local variables
may actually be
stored in the
coroutine state

```
}
```

```
using promise_type =  
    coroutine_traits<lazy<int>>::promise_type;
```

```
{  
    promise_type p;  
    auto r = p.get_return_object();  
    co_await p.initial_suspend();  
    try {  
        int result = hard_work();  
        p.return_value(result);  
        goto final_suspend;  
    } catch (...) {  
        p.unhandled_exception();  
    }  
    final_suspend:  
    co_await p.final_suspend();  
    destroy: ;  
}
```

the return
object is really
like a local
variable

Coroutine state

stores promise, parameter copies, local variables
and any information needed to resume the
coroutine at the correct point

dynamically allocated (can be elided in certain
circumstances)

created when the coroutine is called

destroyed when control flows off the end of the
transformed coroutine body

| |
|------------------|
| promise |
| parameter copies |
| local variables |
| resume point |

co_await transformation

All coroutines `co_await`
at least once, on
`p.initial_suspend()`

```
auto y = co_await x;
```

```
auto h = coroutine_handle<P>::from_promise(p);  
auto a = p.await_transform(x); // optional  
auto e = operator co_await(a); // optional  
if (!e.await_ready()) {  
    // suspended  
    e.await_suspend(h);  
    if (/* still suspended */) {  
        // return to caller  
resume: ;  
    }  
    // resumed  
}  
auto y = e.await_resume();
```

co_yield transformation

co_yield is almost the same as co_await

```
auto y = co_yield x;
```

```
auto h = coroutine_handle<P>::from_promise(p);
auto a = p.yield_value(x);
auto e = operator co_await(a); // optional
if (!e.await_ready()) {
    // suspended
    e.await_suspend(h);
    if (/* still suspended */)
        // return to caller
resume: ;
}
// resumed
}
auto y = e.await_resume();
```

for co_await transformation

```
for co_await (auto x : xs) {
```

```
    use(x);
```

```
}
```

```
{
```

```
    auto __end = xs.end();
```

```
    for (auto __begin = co_await xs.begin();
```

```
        __begin != __end;
```

```
        co_await ++__begin) {
```

```
        auto x = *__begin;
```

```
        use(x);
```

```
    }
```

```
}
```

Coroutine handle

non-owning;
destructor
doesn't do
anything

might not refer to a
coroutine at all

provides access to the
promise

can get one from
the promise

only way to resume or
destroy a coroutine

```
template <typename Promise>
struct coroutine_handle {
    constexpr explicit
        operator bool() const noexcept;

    Promise& promise() const;

    static coroutine_handle
        from_promise(Promise&);

    void resume();
    void destroy();
};
```

Return values

Whenever a coroutine suspends or returns, it returns control to its caller (the first time) or its resumer (subsequent times).

```
lazy<int> v = calculate(); // caller
```

When returning to its caller, the return object is used - it is converted to the return type of the coroutine.

```
h.resume(); // resumer  
h.destroy(); // resumer
```

When returning to a resumer there is no return value.

Implementation for lazy 1/n

```
namespace stdx = std::experimental;
```

```
template <typename T>  
struct lazy_promise;
```

```
template <typename T>  
class lazy {  
private:
```

for brevity;
don't do this in
a header file!

lazy and
lazy_promise
refer to each other

to control the
coroutine

```
using handle_type =  
    stdx::coroutine_handle<  
        lazy_promise<T>>;
```

```
lazy(lazy_promise<T>& p)  
    : h(handle_type::  
        from_promise(p)) {}
```

```
handle_type h;
```

```
friend lazy_promise<T>;
```


Implementation for lazy 2/n

```
public:
```

```
    T get() {  
        if (!h.promise().value) {  
            h.resume();  
        }  
        return *h.promise().value;  
    }
```

executes
remainder of
coroutine
synchronously

```
    ~lazy() {  
        if (h) h.destroy();  
    }  
};
```

also need
copy/move
constructor and
assignment

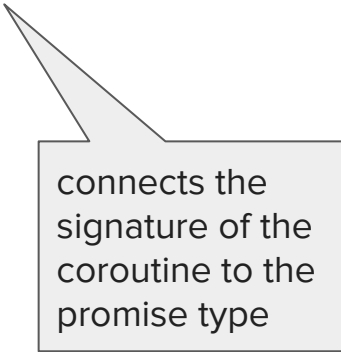
```
template <typename T>  
struct lazy_promise {  
    std::optional<T> value;
```

to store the
lazily-computed
value

```
    auto get_return_object() {  
        return lazy<T>(*this); }  
    auto initial_suspend() {  
        return stdx::suspend_always{}; }  
    void return_value(T x) {  
        value = std::move(x); }  
    void unhandled_exception() {}  
    auto final_suspend() {  
        return stdx::suspend_always{}; }  
};
```

Implementation for lazy 3/3

```
namespace std::experimental {  
    template <typename T, typename... Args>  
    struct coroutine_traits<lazy<T>, Args...> {  
        using promise_type = lazy_promise<T>;  
    };  
}
```



connects the
signature of the
coroutine to the
promise type

I want to use this stuff now!

Working coroutine implementations:

- Visual Studio 2017 (/await)
- clang 5.0 with libc++ 5.0 (-fcoroutines-ts -stdlib=libc++)

Coroutine abstraction libraries:

- cppcoro (<https://github.com/lewissbaker/cppcoro>)
 - task, generator, async_generator, async_mutex, ...
- range-v3 (<https://github.com/ericniebler/range-v3>)
 - generator
- others? Let me know!

Summary

Coroutines are a generalization of functions that allow for many new and interesting control flows.

Usual suspects:

- generators
- `async`
- `async generators`

They require support from library types in order to define what the syntax means.

Thank you!

Some unexpected things you can do

Some operations returning optional

```
optional<string> read_word(istream& s);  
optional<int> parse_int(istream& s);  
optional<double> parse_double(istream& s);
```

```
// “3 1.2 3.4 5.6” -> vector{1.2, 3.4, 5.6}  
optional<vector<double>> parse_vector(istream& s);
```

Composing operations returning optional

```
optional<vector<double>> parse_vector(istream& s) {  
    optional<int> n = parse_int(s);  
    if (!n) return {};  
    vector<double> result;  
    for (int i = 0; i < *n; ++i) {  
        optional<double> x = parse_double(s);  
        if (!x) return {};  
        result.push_back(*x);  
    }  
    return result;  
}
```


Composing operations in a coroutine

```
optional<vector<double>> parse_vector(istream& s) {  
    int n = co_await parse_int(s);  
  
    vector<double> result;  
    for (int i = 0; i < n; ++i) {  
  
        result.push_back(co_await parse_double(s));  
    }  
    co_return result;  
}
```

How can this be?

need to specialize `coroutine_traits` for functions returning `optional`

```
namespace std::experimental {  
    template <typename T, typename... Args>  
    struct coroutine_traits<optional<T>, Args...> { // UB!  
        using promise_type = optional_promise<T>;  
    };  
}
```

- this is not allowed if we're talking about `std::optional` - there needs to be a user-defined type in there somewhere
- but it works anyway

How does it work?

It's all in the awaitable:

```
template <typename T>
struct optional_promise {
    template <typename U>
    auto await_transform(optional<U> e) {
        return optional_awaitable<U>{move(e)};
    }
};
```

It's all in the awaitable

```
template <typename U>
struct optional_waitable {
    optional<U> o;
```

suspend if the
optional does not
contain a value

```
    auto await_ready() { return o.has_value(); }
    auto await_resume() { return o.value(); }
```

if it does contain
a value, return
that value

```
template <typename T>
void await_suspend(coroutine_handle<maybe_promise<T>> h) {
    h.promise().data->emplace(nullopt);
    h.destroy();
}
};
```

if no value,
destroy the
coroutine

But I want more information about what failed

Luckily for you, the exact same approach works with types like `expected`.

```
enum class ParseError { EndOfStream, BadFormat };
```

```
template <typename Value>
```

```
using parse_result = expected<Value, ParseError>;
```

```
parse_result<string> read_word(istream& s);
```

```
parse_result<int> parse_int(istream& s);
```

```
parse_result<double> parse_double(istream& s);
```

```
parse_result<vector<double>> parse_vector(istream& s);
```

I want to play with optional coroutines

- https://github.com/toby-allsopp/coroutine_monad
 - works with `std::optional`
 - but only on clang
 - has a basic `expected` implementation
- `folly::Optional` (<https://github.com/facebook/folly>)
 - not `std::optional`
 - works on clang *and* MSVC
 - would be pretty easy to get `folly::Expected` working as well

Type-erased callable

- This idea is from Gor Nishanov
- The observation is that the coroutine body can hold onto more type information than is reflected in the return type of the coroutine

Type-erased callable

```
template <typename Ret,  
          typename... Args>  
class func {  
    template <typename F>  
    static func create(F f) {  
        co_yield f;  
    }  
  
    Ret operator()(Args... args) {  
        h.promise().args = {args...};  
        h.resume();  
        return h.promise().ret;  
    }  
};
```

function type
erased here

function type
recovered
here

```
template <typename Ret,  
          typename... Args>  
struct func_promise() {  
    tuple<Args...> args;  
    Ret ret;  
  
    template <typename F>  
    void yield_value(F f) {  
        ret = apply(f, args);  
    }  
};
```


Type-erased callable

- If the compiler can see what you put in and where you call it then all of the coroutine machinery can be optimized away
- Otherwise, overhead is probably similar to `std::function`
- An intriguing technique
- Some code to start playing with:
https://github.com/toby-allsopp/coroutine_func